

Summary and resampling statistics

Spyros Samothrakis
Research Fellow, IADS
University of Essex

January 24, 2017

About

Summary statistics

Confidence Intervals

Hypothesis testing (A/B testing)

Conclusion

SUMMARY STATISTICS AND RESAMPLING STATISTICS

- ▶ Today we are going to discuss summary statistics and resampling statistics
 - ▶ Summary statistics try to capture the “essence” of a set of observations (the sample)
 - ▶ Resampling statistics create different samples from the original sample in order to gain further insights
- ▶ Resampling statistics are far more intuitive to understand than using t-tests (I think...)

AN EXAMPLE PROBLEM

- ▶ Let's say that a journalist was tasked with finding the salaries of a business
- ▶ But could only find through friends and acquaintances the salaries of certain employees

| Employee ID | Salary |
|-------------|--------|
| 1 | 10000 |
| 2 | 100000 |
| 3 | 200000 |
| 4 | 140000 |
| 5 | 12000 |
| 6 | 13000 |
| 7 | 140000 |
| 8 | 15000 |
| 9 | 120000 |

(CONTINUED TABLE)


| Employee ID | Salary |
|-------------|--------|
| 10 | 11000 |
| 11 | 8000 |
| 12 | 9000 |
| 13 | 14000 |
| 14 | 14000 |
| 15 | 5000 |
| 16 | 18000 |
| 17 | 6000 |
| 18 | 18000 |
| 19 | 15000 |
| 20 | 19000 |
| 21 | 12000 |

LET'S PLOT

```
df = pd.read_csv('./customers.csv')
```


```
# There are far  
# better ways of doing this  
data = df.values.T[1]
```

```
sns_plot = sns.distplot(df,  
    bins=20,  
    kde=False,  
    rug=True).get_figure()
```



HISTOGRAM PLOT

```
sns_plot2 = sns.distplot(data,  
    bins=20,  
    kde=False,  
    rug=True).get_figure()
```



MEASURES OF CENTRAL TENDENCY

- ▶ (Sample) mean

- ▶ $\mu = \frac{1}{N} \sum_{i=1}^N X_i$

- ▶ (Sample) median

- ▶ Rank X_i

- ▶
$$M = \begin{cases} X_{N/2+1} & \text{if } n \text{ is odd} \\ (X_{N/2} + X_{(N+1)/2})/2 & \text{if } n \text{ is even} \end{cases}$$

- ▶ In the salary data

- ▶ $\mu = 42809.523810$

- ▶ $M = 14000.000000$

MEASUREMENTS OF DISPERSION

- ▶ (Sample) Standard deviation

- ▶ $\sigma = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (X_i - \bar{X})^2}$

- ▶ Variance is σ^2

- ▶ Median absolute deviation

- ▶ $MAD = M(|X_i - M(X)|)$

- ▶ In our data we have:

- ▶ $\sigma^2 = 3230916099.773242$



- ▶ $\sigma = 56841.147946$

- ▶ $MAD : 4000.000000$

SALES

- ▶ A company has recorded their sales for 14 days
- ▶ They want to understand their data
- ▶ Let's plot

HISTOGRAM PLOT OF SALES



SUMMARY STATISTICS

$$\mu = 9.214$$

$$M : 8.500000$$


$$\sigma^2 : 32.311$$

$$\sigma : 5.684296$$


$$M : 2.500$$

Note that there are tons of other summary statistics, this is practically for illustration purposes only


NORMAL DISTRIBUTION




UNIFORM DISTRIBUTION



NORMAL HIGH VARIANCE DISTRIBUTION



PARETO DISTRIBUTION



ARE WE CONFIDENT WE GOT THE RIGHT MEAN?

- ▶ How confident should the journalist or the analyst be about their summary statistics?
- ▶ If they sampled another 14 days, maybe the sale numbers would be completely different?
- ▶ We would like to build some notion of “confidence intervals” (CI)
 - ▶ Get a measure of “If I do this sampling process over and over again, what would I expect to be seeing?”
- ▶ We are going to take the above statement seriously
 - ▶ And introduce the bootstrap!

THE BOOTSTRAP

- ▶ We are going to use a method called the bootstrap to create those CIs
- ▶ Very popular, computational method
- ▶ DiCiccio, Thomas J., and Bradley Efron. “Bootstrap confidence intervals.” Statistical science (1996): 189-212.
- ▶ You will see this name (bootstrap) used quite often in scientific contexts
 - ▶ It refers to a self-starting process
 - ▶ The mind “understanding itself”
 - ▶ Pulling yourself up by the bootstraps
- ▶ Hard to do without a machine

BOOTSTRAPPING (1)

- ▶ Ideally, we could possibly sample again and again from the population
 - ▶ i.e. the journalist would go over to a different set of friends
 - ▶ Ask them to get her some salaries
 - ▶ Repeat
- ▶ Once we have a collection of different means we can say that a mean will fall within a certain range with a certain probability
 - ▶ But this is almost impossible
- ▶ We can use our sample however in a smart way
 - ▶ Resample from the sample!

BOOTSTRAPPING (2)

- ▶ Sample with replacement from the data you have already
 - ▶ Create $\{1 \dots B\}$ bootstraps of the same size
 - ▶ Let's assume each observation in the initial dataset is X_i , where i is the order appearing

$$X^1 = X_4^1, X_5^1, X_3^1, X_5^1 \dots$$

$$X^2 = X_3^2, X_7^2, X_7^2, X_8^2 \dots$$

$$X^{\dots} = \dots$$

$$X^B = X_8^B, X_3^B, X_2^4, X_4^1 \dots$$

BOOTSTRAPPING (3)

- ▶ Let's do one one example
- ▶ $X = \{1,0,1,2\}$
- ▶ Let's draw three samples
 - ▶ I will simulate the dice rolls


BOOTSTRAPPING (4)

- ▶ Get the mean for each sample (since this is what we are interested in)
- ▶ We can now rank the means
- ▶ We remove the bottom 10% and the top 10% to find $\gamma = 0.80$
- ▶ For the salary data


$X = [6.86, 7.29, 7.86, 8.14$
 $8.36, 8.79, 8.86, 9.14$
 $9.29, 9.5, 9.5, 9.71$
 $10.36, 11.14, 11.14, 13.21]$

- ▶ What about if I was interested in $\gamma = 0.90$?
- ▶ What about if I was interested in $\gamma = 0.95$?

SALARIES



SALES



WHAT CAN WE SAY ABOUT THE MEANS NOW?

- ▶ Salaries mean is...
- ▶ Sales mean is...
- ▶ We can do bootstrap to estimate *any* quantity we want as long as the distribution has a defined variance and mean
 - ▶ i.e. not always
- ▶ But for most practical matters, yes

DATA BIAS

- ▶ I have described a very biased process of collecting samples
 - ▶ The journalist asked her friends
 - ▶ All her friends love football
 - ▶ What he might actually have learned is the salary of football loving employees
- ▶ How about the sales figures?
 - ▶ Was there anything extra-ordinary on the day these measurements were taken?
 - ▶ Maybe it was Christmas
- ▶ Be very careful to randomise properly, and if not at least take care to state your bias

A/B TESTING

- ▶ Suppose you had two versions of a website
 - ▶ and you would like to check if the newer version is better
- ▶ Two versions of an e-mail
 - ▶ and you would like to check if the newer, fancier version is better
- ▶ A new drug
 - ▶ and you would like to see if it actually cures
- ▶ A zombie apocalypse
 - ▶ and you have found a serum to cure zombiness

HYPOTHESIS TESTING

- ▶ Same as A/B testing
- ▶ Not just limited to binary cases
- ▶ The name people used to call the same procedure when testing for
 - ▶ Drug effects
 - ▶ Physical effects
 - ▶ Quality management
- ▶ A lot of Data science concepts are just “re-imaginings”

EXAMPLE PROBLEM

- ▶ A company sends out e-mails
 - ▶ Various promotions and news content
 - ▶ They want users to click on the links and get on their website
 - ▶ They already have an e-mail format
 - ▶ Mark from marketing comes up with an e-mail with improved content
- ▶ Is it better?
 - ▶ Without causing too much disruption

HYPOTHESIS TESTING

- ▶ They send 11 e-mails of the usual type (control)
- ▶ They also send 11 e-mails of the new design (test)

```
old = np.array([0,1,1,1,0,1,1,0,0,1,0])
```

```
new = np.array([0,1,1,0,1,1,0,1,1,1,0])
```

$$\mu_{old} = 0.18$$

$$\mu_{new} = 0.455$$

$$t_{obs} = \mu_{new} - \mu_{old} = 0.27$$

Should they change?

HYPOTHESIS FORMING

H_0 : The two e-mails have no difference (their means are equal) - this is called the *null* hypothesis

H_1 : The second e-mail is better, and thus has a higher mean

- ▶ Set $\alpha = 0.05$, or equivalently the 95% CI t_{obs} does not contain H_1
- ▶ The CI of H_0 does not contain H_1
- ▶ What is the probability of observing something as extreme as we just observed by pure chance?

PERMUTATION TESTING (1)

- Merge all the data into a new array

```
array([0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0,  
       0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0])
```

- Permute it random, i.e. form a new array from the same elements

```
array([0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1,  
       0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0])
```


PERMUTATION TESTING (2)

- ▶ Split again into new and old


```
pold = np.array([0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1])  
pnew = np.array([0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0])
```

- ▶ Record if the value of the test was more extreme or not
 - ▶ $t_{perm} = \mu_{pnew} - \mu_{pold}$
 - ▶ $t_{perm} > t_{obs}$
- ▶ Keep on permuting and recording
- ▶ Find the number of times $t_{perm} > t_{obs}$
 - ▶ Divide by the number of permutations you used
- ▶ You call that number your *p-value*

PERMUTATION TESTS (3)

- ▶ If you do this for 19,000 permutations you get $p - value = 0.032$
- ▶ Hence we can conclude 5 out of a 100 times you will get a higher difference in means
- ▶ Find out if this number is smaller than $\alpha = 0.05$
- ▶ If yes, you can reject the H_0 (which it is)

PERMUTATION TEST (4)



ANOTHER EXPERIMENT

- ▶ Bob decides that adding a sound to the e-mail should increase user clicking even more
- ▶ Thinking that it his solution is better for sure, he sends more e-mails with sounds (i.e. the new version)
 - ▶ Not exactly A/B testing, but he seems eager...
- ▶ Results come back and he had to somehow show that his new e-mail procedure is better

SOME DATA ANALYSIS

```
old = np.array([0,1,1,1,0,1,1,0,0,1,0])  
new = np.array([0,1,1,0,1,1,0,1,1,1,0,0,1,1,1,1,1,1,1])
```

$$\mu_{old} = 0.546$$

$$\mu_{new} = 0.73$$

$$t_{obs} = \mu_{new} - \mu_{old} = 0.19$$

RESULTS

- ▶ With 19,000 permutations we get a $p = 0.07$
- ▶ Thus we have failed to reject the null hypothesis
- ▶ Does not mean that the sound doesn't have any impact
- ▶ Just that we can't tell the impact

ERRORS

- ▶ Type I error: rejecting H_0 even though it is true
- ▶ Type II error: failing to reject H_0 even though it is false

| | H_0 is true | H_0 is false |
|----------------------|-------------------------------|--------------------------------|
| Reject H_0 | Type I error (false positive) | Correct Inference |
| Fail to reject H_0 | Correct inference | Type II error (false negative) |

SPECIFICITY

- ▶ False positive rate refers to the level we set α
- ▶ $1 - \alpha$ is the *specificity* of the test, the proportion of true negatives
- ▶ The higher, the more susceptible the test is to Type I errors
- ▶ Think of this as raising false alarms

SENSITIVITY

- ▶ False negative rate refers to another parameter, which we haven't set at all for now, called β
- ▶ $1 - \beta$ is the *sensitivity* or power of a test / the ratio of true positives
- ▶ The higher it is, the more we are bound to do Type II errors
- ▶ Think of this as failure to detect a phenomenon
- ▶ It is indirectly influenced by effect size and sample size
- ▶ “Surely you only need one of them!” (No!)

POWER ANALYSIS

- ▶ A question that would naturally rise up is how many samples do we need to collect, if we are to perform a study within a certain error
- ▶ No easy solution
- ▶ In practice, sample as much as you can
- ▶ See previous studies in the literature
- ▶ If you have done a study before, use the bootstrap!
 - ▶ How?
- ▶ You might be tempted to increase α , but this will increase your chance for a Type I error

A MORE “HACKISH IDEA”

- ▶ Get the confidence intervals for both populations
- ▶ If they overlap, fail to reject H_0
- ▶ If not, reject H_0
- ▶ Very tempting to do this
 - ▶ Actually you can
 - ▶ It's a bit more conservative, but people do it all the time
 - ▶ Not thaaaaat bad if the samples are independent

Schenker, Nathaniel, and Jane F. Gentleman. “On judging the significance of differences by examining the overlap between confidence intervals.” *The American Statistician* 55.3 (2001): 182-186.

P-HACKING

“In the course of collecting and analyzing data, researchers have many decisions to make: Should more data be collected? Should some observations be excluded? Which conditions should be combined and which ones compared? Which control variables should be considered? Should specific measures be combined or transformed or both?”

Simmons, Joseph P., Leif D. Nelson, and Uri Simonsohn.

“False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant.”

Psychological science 22.11 (2011): 1359-1366.

CONCLUDING

- ▶ Hypothesis testing is used quite extensively
- ▶ And abused more often
- ▶ Cross validation?
- ▶ Real life problems (usually) have more data and are more noisy
 - ▶ But you can send e-mails, get clicks etc. trivially
- ▶ If there is one thing to keep from this lecture is the use of bootstrapping to learn parameter confidence intervals
 - ▶ We will use bootstrap later on this module when we are going to model things