

| ABOUT | COLLABORATIVE FILTERING | CONTENT-BASED FILTERING | IMPLICIT FEEDBACK | CONCLUSION |
|-------|-------------------------|-------------------------|-------------------|------------|
|-------|-------------------------|-------------------------|-------------------|------------|

# Recommender Systems

Spyros Samothrakis  
Research Fellow, IADS  
University of Essex

February 14, 2017

1 / 38

| ABOUT | COLLABORATIVE FILTERING | CONTENT-BASED FILTERING | IMPLICIT FEEDBACK | CONCLUSION |
|-------|-------------------------|-------------------------|-------------------|------------|
|-------|-------------------------|-------------------------|-------------------|------------|

About

Collaborative filtering

Content-based filtering

Implicit feedback

Conclusion

2 / 38

| ABOUT | COLLABORATIVE FILTERING | CONTENT-BASED FILTERING | IMPLICIT FEEDBACK | CONCLUSION |
|-------|-------------------------|-------------------------|-------------------|------------|
|-------|-------------------------|-------------------------|-------------------|------------|

## RECOMMENDER SYSTEMS

- ▶ We will discuss one of the most popular applications of data science
  - ▶ Recommender Systems
- ▶ Every website does it
- ▶ Recommender Systems match users with items
- ▶ Users under constant information overload
- ▶ Think songs, foods, drinks, movies, news

3 / 38

| ABOUT | COLLABORATIVE FILTERING | CONTENT-BASED FILTERING | IMPLICIT FEEDBACK | CONCLUSION |
|-------|-------------------------|-------------------------|-------------------|------------|
|-------|-------------------------|-------------------------|-------------------|------------|

## EXAMPLES

- ▶ This is even done offline in the service industry!
- ▶ Can you think of other examples?

4 / 38

## COLLABORATIVE FILTERING

- ▶ Collaborative filtering is an effort to predict how products/items will be rated by a user, given previous ratings (from both the user and others)
- ▶ This prediction can help us recommend to the user only items that we think she will rate highly
- ▶ Latent Factor Models - Netflix Challenge (1M\$)- Simon Funk

5 / 38

## SAME SAMPLE DATA

|   | The Call of Cthulhu | Frankenstein | Dracula | Neuromancer | Dune |
|---|---------------------|--------------|---------|-------------|------|
| 0 | 6                   | 0            | 0       | 7           | NaN  |
| 1 | 5                   | 0            | 5       | 6           | 5    |
| 2 | 9                   | NaN          | 4       | NaN         | 8    |
| 3 | 4                   | NaN          | 2       | 5           | 6    |
| 4 | 4                   | NaN          | 4       | 6           | 0    |
| 5 | 6                   | 3            | 8       | 5           | 7    |
| 6 | NaN                 | 6            | NaN     | 6           | 7    |
| 7 | NaN                 | 1            | 1       | 6           | NaN  |
| 8 | NaN                 | NaN          | 2       | NaN         | 9    |
| 9 | NaN                 | 3            | 4       | 5           | 7    |


6 / 38

## FACTORS

- ▶ We are going to base our predictions on “hidden” qualities of the items
- ▶ For example, food can have different levels of spiciness, a drink different levels of bitterness
- ▶ We term these qualities “factors”
- ▶ A sensible way of describing items would be to see them as a collection of “factors”
  - ▶ But our data is just ratings!
- ▶ Thus, our factors are “latent”, i.e. hidden!

7 / 38

## EXAMPLE FACTORS FOR OUR DATA



8 / 38

## FACTORS AND USER PREFERENCES

- ▶ Let's assume  $n$  factors
- ▶ We can encode factors as a real valued vector  
 $\text{item\_factors}_i = [\phi_0, \phi_1, \dots, \phi_{n-1}]$
- ▶ For example “The Call of the Cthulhu” can be encoded as  
 $\text{item\_factors}_0 = [3, 8]$
- ▶ Each user now can have preferences over factors, encoded as weights  $\text{user\_preferences}_j = [w_0, w_1, \dots, w_{n-1}]$
- ▶ The weight vector contains user preferences, e.g.  
 $\text{user\_preferences}_0 = [0.5, 0.8]$
- ▶ But we don't have any user weights nor any item factors - generate some random!

9 / 38

## SOME RANDOM DATA

- ▶ Each row in *user\_preferences* represents the preferences of a user, while each row in *item\_factors* represents the factors of an item

```
array([[ 0.092,  0.783],
       [ 0.78 ,  0.488],
       [ 0.844,  0.062],
       [ 0.68 ,  0.549],
       [ 0.212,  0.43 ],
       [ 0.961,  0.023],
       [ 0.659,  0.31 ],
       [ 0.92 ,  0.769],
       [ 0.817,  0.452],
       [ 0.834,  0.887]])
array([[ 0.338,  0.519],
       [ 0.69 ,  0.256],
       [ 0.363,  0.93 ],
       [ 0.004,  0.112],
       [ 0.608,  0.104]])
```

$\text{rating}[0][0] \leftarrow 0.437 = 0.092 * 0.338 + 0.783 * 0.519$

Far away from the observed rating of 6

10 / 38

## PREDICT RATINGS

- ▶ Example python code
- ▶ Still random values...

```
def predict_rating(user_row, item_row):
    """ Predict a rating given a user_row and an item_row.
    """
    user_values = user_preferences[user_row]
    item_values = item_factors[item_row]
    return user_values.dot(item_values)
```

11 / 38

## TRAINING FOR A SINGLE EXAMPLE

- ▶ *user\_preferences* and *item\_factors* have random values!
- ▶ Find the difference between the real and the predicted rating (“how far away am I from the goal?”)
- ▶ Multiply by small learning rate  $\alpha = 0.0001$  (“Don't take my measurement so seriously”)
- ▶ Move *user\_preferences* and *item\_factors* towards the correct value, following the negative of the gradient (“Let's move towards the direction of the most abrupt change”)

12 / 38

## TRAINING CODE

```
def train(user_row,item_row,rating,alpha = 0.0001):
    """ Adapt the values of user_preferences and item_factors
        to match the ones predicted by the users
    """
    err = alpha * (rating - predict_rating(user_row, item_row))
    user_preferences[user_row] += err * item_factors[item_row]
    item_factors[item_row] += err * user_preferences[user_row]
```

13 / 38

## TRAINING USING ALL DATA

- ▶ “Latent Factors” because we have never really observed them, we can only infer them
- ▶ Loop over all *user\_preferences* and *item\_factors*
- ▶ Ignore cells with no value (“NaN” cells)
- ▶ Repeat until some criterion (in our case, 30,000 iterations)

14 / 38

## TRAINING USING ALL DATA CODE

```
def sgd_svd(iterations = 30000):
    """ Iterate over all users and all items and train for
        a certain number of iterations
    """
    for i in range(0,iterations):
        for user_row in range(0,user_preferences.shape[0]):
            for item_row in range(0,item_factors.shape[0]):
                rating = user_ratings[user_row][item_row]
                if(not np.isnan(rating)):
                    train(user_row,item_row,rating)
```

15 / 38

## RECONSTRUCTING DATA / PREDICTING UNSEEN RATINGS

- ▶ Run *sgd\_svd()* and print the updated tables

```
array([[ 1.705,  0.486],
       [ 1.857,  0.484],
       [ 2.373,  1.311],
       [ 0.988,  1.495],
       [ 2.519, -2.088],
       [ 1.868,  1.235],
       [ 1.367,  1.801],
       [ 1.405,  0.628],
       [ 0.133,  3.453],
       [ 1.562,  1.235]])
array([[ 2.713,  1.266],
       [-1.125,  4.09 ],
       [ 1.847,  0.514],
       [ 3.09 ,  0.807],
       [ 2.079,  2.522]])
```

$rating[0][0] \leftarrow 1.705 * 2.713 + 0.486 * 1.266 \approx 5.2$ , not 6, but close!

16 / 38

## VISUAL COMPARISON

- Calculate all predicted values and pretty print

|   | The Call of Cthulhu | Frankenstein   | Dracula       | Neuromancer   | Dune          |
|---|---------------------|----------------|---------------|---------------|---------------|
| 0 | (6.000 5.209)       | (0.000 0.029)  | (0.000 3.397) | (7.000 5.699) | (nan 4.587)   |
| 1 | (5.000 5.662)       | (0.000 -0.046) | (5.000 3.699) | (6.000 6.208) | (5.000 4.956) |
| 2 | (9.000 8.112)       | (nan 2.705)    | (4.000 5.098) | (nan 8.443)   | (8.000 8.198) |
| 3 | (4.000 4.564)       | (nan 4.518)    | (2.000 2.651) | (5.000 4.262) | (6.000 5.799) |
| 4 | (4.000 4.214)       | (nan -9.302)   | (4.000 3.425) | (6.000 6.127) | (0.000 0.017) |
| 5 | (6.000 6.617)       | (3.000 3.004)  | (8.000 4.100) | (5.000 6.756) | (7.000 7.003) |
| 6 | (nan 5.960)         | (6.000 5.735)  | (nan 3.473)   | (6.000 5.593) | (7.000 7.508) |
| 7 | (nan 4.577)         | (1.000 0.950)  | (1.000 2.918) | (6.000 4.858) | (nan 4.397)   |
| 8 | (nan 4.494)         | (nan 12.716)   | (2.000 2.010) | (nan 2.851)   | (9.000 8.985) |
| 9 | (nan 5.769)         | (3.000 3.336)  | (4.000 3.523) | (5.000 5.774) | (7.000 6.389) |

- For user 2, recommend “Neuromancer” and ignore “Frankenstein”
- Observe how reconstruction is not perfect - multiple reasons (e.g. data shuffling? mini-batches? more factors? more training)

17 / 38

## OTHER OUTCOME SIGNALS

- We have used ratings
- But this is not the only possible outcome
- One can use other signals as well
  - User have searched for certain films
  - Users have clicked on certain films

18 / 38

## HOW ABOUT ADDING KNOWN FACTORS/FEATURES?

- Until now we only had latent factors
- But latent factors arise only when you actually have some a good number of \$ pairs for a user
  - What if the user just joined the website?
- What if you don't have any?
  - Or what if you have further observations about a user

19 / 38

## ITEM DESCRIPTIONS

- Instead of based only on latent factors, we can base our predictions on known observations
- For example each book can have:
  - Genre
  - Data published
  - Age of intended audience
  - Author
- Each film can have
  - Genre
  - Director
  - Age of intended audience

20 / 38

## USER DESCRIPTIONS

- ▶ But we might have data collected about a user as well
  - ▶ Age
  - ▶ Sex
  - ▶ Country of birth
  - ▶ Native language
- ▶ All kinds of data

21 / 38

## EXPLICIT KNOWLEDGE

- ▶ You ask the user questions explicitly
  - ▶ What kind of books do you like?
  - ▶ What is the maximum price you would pay for a book?

22 / 38

## CONTENT BASED FILTERING

- ▶ One can use all features defined on the user and the item
- ▶ Create a classifier and do the predictions using the classifier
- ▶ We have very few samples
  - ▶ So we are going to use a linear classifier

23 / 38

## ITEM FEATURE MATRIX

Feature 0: First critic score of the book

Feature 1: Second critic score of the book

|   | Critic0 | Critic1 |
|---|---------|---------|
| 0 | 0.3     | 0.9     |
| 1 | 0.9     | 0.3     |
| 2 | 0.6     | 0.4     |
| 3 | 0.2     | 0.1     |
| 4 | 0.7     | 0.8     |
| 5 | 0.9     | 0.1     |

24 / 38

## USER FEATURE MATRIX

Feature 0: Male/Female

Feature 1: Over 60

|   | Sex | Over60 |
|---|-----|--------|
| 0 | 1.0 | 0.0    |
| 1 | 0.0 | 1.0    |
| 2 | 0.0 | 0.0    |
| 3 | 1.0 | 0.0    |
| 4 | 0.0 | 1.0    |
| 5 | 0.0 | 0.0    |
| 6 | 0.0 | 0.0    |
| 7 | 1.0 | 0.0    |
| 8 | 0.0 | 1.0    |
| 9 | 1.0 | 0.0    |

25 / 38

## COMBINING THE FEATURES

- We need to build a set of features for training for each person/item combo

|   | Sex | Over60 | key | user_id | Critic0 | Critic1 | item_id | rating |
|---|-----|--------|-----|---------|---------|---------|---------|--------|
| 0 | 1.0 | 0.0    | 0   | 0       | 0.3     | 0.9     | 0       | 8.0    |
| 1 | 1.0 | 0.0    | 0   | 0       | 0.9     | 0.3     | 1       | 2.0    |
| 3 | 1.0 | 0.0    | 0   | 0       | 0.2     | 0.1     | 3       | 5.0    |
| 4 | 1.0 | 0.0    | 0   | 0       | 0.7     | 0.8     | 4       | 4.0    |
| 0 | 0.0 | 1.0    | 0   | 1       | 0.3     | 0.9     | 0       | 3.0    |
| 1 | 0.0 | 1.0    | 0   | 1       | 0.9     | 0.3     | 1       | 2.0    |
| 3 | 0.0 | 1.0    | 0   | 1       | 0.2     | 0.1     | 3       | 7.0    |
| 4 | 0.0 | 1.0    | 0   | 1       | 0.7     | 0.8     | 4       | 7.0    |
| 0 | 0.0 | 0.0    | 0   | 2       | 0.3     | 0.9     | 0       | 9.0    |
| 2 | 0.0 | 0.0    | 0   | 2       | 0.6     | 0.4     | 2       | 7.0    |
| 3 | 0.0 | 0.0    | 0   | 2       | 0.2     | 0.1     | 3       | 8.0    |
| 4 | 0.0 | 0.0    | 0   | 2       | 0.7     | 0.8     | 4       | 5.0    |
| 2 | 1.0 | 0.0    | 0   | 3       | 0.6     | 0.4     | 2       | 7.0    |
| 3 | 1.0 | 0.0    | 0   | 3       | 0.2     | 0.1     | 3       | 8.0    |
| 4 | 1.0 | 0.0    | 0   | 3       | 0.7     | 0.8     | 4       | 9.0    |
| 1 | 0.0 | 1.0    | 0   | 4       | 0.9     | 0.3     | 1       | 1.0    |
| 2 | 0.0 | 1.0    | 0   | 4       | 0.6     | 0.4     | 2       | 8.0    |
| 3 | 0.0 | 1.0    | 0   | 4       | 0.2     | 0.1     | 3       | 3.0    |

26 / 38

## TEST SET

- We are looking to predict this:

|   | Sex | Over60 | key | user_id | Critic0 | Critic1 | item_id | rating |
|---|-----|--------|-----|---------|---------|---------|---------|--------|
| 2 | 1.0 | 0.0    | 0   | 0       | 0.6     | 0.4     | 2       | NaN    |
| 2 | 0.0 | 1.0    | 0   | 1       | 0.6     | 0.4     | 2       | NaN    |
| 1 | 0.0 | 0.0    | 0   | 2       | 0.9     | 0.3     | 1       | NaN    |
| 0 | 1.0 | 0.0    | 0   | 3       | 0.3     | 0.9     | 0       | NaN    |
| 1 | 1.0 | 0.0    | 0   | 3       | 0.9     | 0.3     | 1       | NaN    |
| 0 | 0.0 | 1.0    | 0   | 4       | 0.3     | 0.9     | 0       | NaN    |
| 3 | 0.0 | 0.0    | 0   | 5       | 0.2     | 0.1     | 3       | NaN    |
| 4 | 0.0 | 0.0    | 0   | 5       | 0.7     | 0.8     | 4       | NaN    |
| 2 | 0.0 | 0.0    | 0   | 6       | 0.6     | 0.4     | 2       | NaN    |
| 2 | 0.0 | 1.0    | 0   | 8       | 0.6     | 0.4     | 2       | NaN    |
| 1 | 1.0 | 0.0    | 0   | 9       | 0.9     | 0.3     | 1       | NaN    |

27 / 38

## CODE - PANDAS MAGIC!!!

```

user_ratings_df = pd.read_csv("user_ratings.csv")
user_features_df = pd.read_csv("user_features.csv")
item_features_df = pd.read_csv("item_features.csv")

user_features_df["key"] = 0
user_features_df["user_id"] = range(0,user_features_df.shape[0])
item_features_df["key"] = 0
item_features_df["item_id"] = range(0,item_features_df.shape[0])

merged_df = pd.merge(user_features_df, item_features_df,left_index=True,on="key")

merged_df["rating"] = map(lambda ids: user_ratings_df.values[ids[1]][ids[2]],
                          merged[["user_id", "item_id"]].itertuples())

train = merged_df.dropna()

test = merged_df[merged.isnull().any(axis=1)]

```

28 / 38

## HYBRID SYSTEMS

- Can we merge the two approaches?
  - Of course we can - various ways of merging
- We will just augment collaborative filtering with standard features for now
- We can add the features as parts of the variables we are going to learn

29 / 38

## REGULARISATION

- We will now add a penalty to features depending on what they are
- Penalty strong for latent features
- But could be the other way around
- Weight decays /  $l_2$  regulariser
  - $w_i = w_i - \alpha(y_p - y)(\phi_i + \lambda w_i)$

30 / 38

## CODE - PREDICT

```
def predict_rating(user_id,item_id):
    """ Predict a rating given a user_id and an item_id.
    """
    user_preference = latent_user_preferences[user_id]
    item_preference = latent_item_features[item_id]

    user_score = user_features_weights[user_id].dot(user_features[user_id])
    item_score = item_features_weights[item_id].dot(item_features[item_id])

    return user_preference.dot(item_preference) + user_score + item_score
```

31 / 38

## CODE - TRAIN

```
def train(user_id, item_id, rating,alpha = 0.0001,
          latent_feature_weight_decay = 0.1,
          user_weight_decay = 0.01,
          item_weight_decay = 0.001):

    prediction_rating = predict_rating(user_id, item_id)
    err = ( prediction_rating - rating );

    user_pref_values = latent_user_preferences[user_id][:]
    latent_user_preferences[user_id] -= alpha *
        err * ( latent_item_features[item_id] +
                latent_feature_weight_decay*
                latent_user_preferences[user_id])

    latent_item_features[item_id] -= alpha *
        err * ( user_pref_values +
                latent_feature_weight_decay*latent_item_features[item_id])

    user_features_weights[user_id] -=alpha * err *( user_features[user_id] +
        user_weight_decay* user_features_weights[user_id])
    item_features_weights[item_id] -=alpha * err * ( item_features_weights[item_id] +
        item_weight_decay* item_features_weights[item_id])

    return err
```

32 / 38



## REMOVING LATENT FACTORS

- ▶ What will happen to the example above if we remove all latent factors?
  - ▶ ... and base our decisions only in known features
- ▶ Factors vs features

33 / 38

## EVALUATING

- ▶ Cross-validation again!
- ▶ We have used all the data in our examples
  - ▶ Is this correct?
- ▶ What would be the proper way of evaluating the system?

34 / 38

## IMPLICIT FEEDBACK

- ▶ How about cases where you
- ▶ When you are recommending books or food user preferences user likes/dislikes something
- ▶ How about if you are recommending
  - ▶ News?
  - ▶ Adverts?
- ▶ You are limited to clicks!
- ▶ But they are not proper feedback
  - ▶ User might not have seen something

35 / 38

## NO NEGATIVE FEEDBACK REGIMES

- ▶ Not clicking is not negative feedback
- ▶ Create a preference matrix - 1 if user has clicked, 0 if she hasn't
- ▶ Weight the importance of each example by  $c(i, j) = 1 + \alpha r(i, j)$ 
  - ▶  $r(i, u)$  is the number of clicks

36 / 38

## CONTEXTUAL BANDITS

- ▶ You can personalise the above scenario even further
- ▶ Send the user some random examples (e.g. news)
  - ▶ With let's say 0.2 probability
- ▶ Seems familiar?
- ▶ Obviously better solutions than  $\epsilon$ -greedy

## CONCLUSION

- ▶ We have seen various instances of recommender systems
- ▶ Again, there are far more complex models
- ▶ But the examples here should have given you a good view about what is out there
- ▶ Also, pandas
  - ▶ Bring your data to a format that is usable by relevant libraries!