

Images, text, video, sound and generative models

Spyros Samothrakis
Research Fellow, IADS
University of Essex

February 7, 2017

About

Video data

Text

Generating data

Conclusion

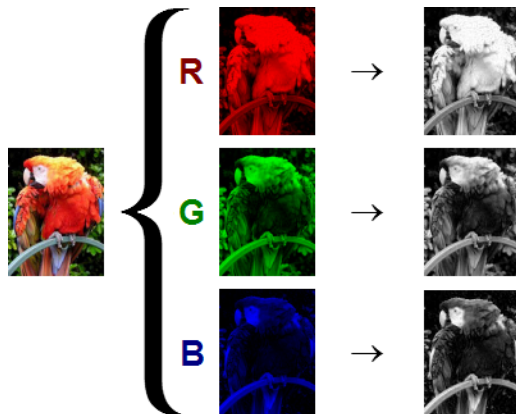
ABOUT

- ▶ We will now turn our attention on data that has less clear structure
- ▶ Sometimes called *unstructured data*
 - ▶ VS *structured* data, i.e. database like tables
- ▶ Is there anything special about this data?
 - ▶ It's the default data humans perceive and generate!
- ▶ Most machine learning benchmarks are on text or image datasets
- ▶ Neural networks excel, but there are other approaches

IMAGE DATA

- ▶ Each image is composed of a number of pixels
- ▶ Pixels have different intensities
- ▶ Also, three channels (RGB)
- ▶ So in effect, we have a three dimensional structure
- ▶ Width x Height x Channels x Intensity
- ▶ 32-bit floating point numbers

RGB EXAMPLE



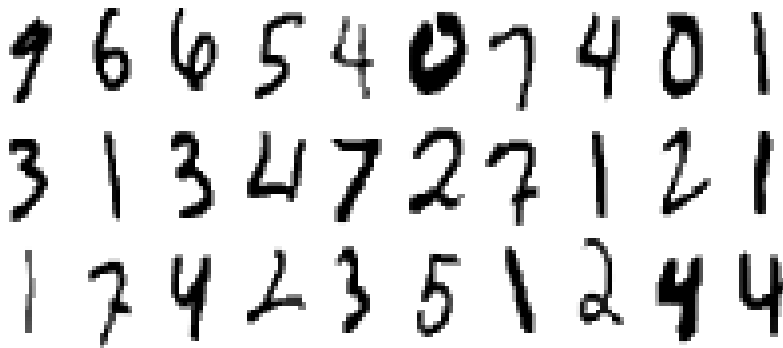
1

¹<http://triplelift.com/2013/07/02/the-complexity-of-image-analysis-part-2-colors/>

MNIST

Very popular benchmark

60,000 training examples, 10,000 test examples, 256 different pixel values, 10 digits,



COMMON IMAGE PREPROCESSING STEPS

- ▶ $28 \times 28 = 784$ features
- ▶ Naive solution
 - ▶ Throw the features to a classifier/regressor
 - ▶ Subtract the mean, divide by the standard deviation
 - ▶ fit/predict
- ▶ This might not work that well

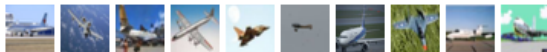
DATA TRUMPS ALGORITHMS

- ▶ It is often tempting to try to find a better algorithm to solve a certain problem
- ▶ But it has been shown time and time again that one much better off by adding more data
- ▶ Problems with neat solutions are very rare, more data
- ▶ *Physics envy*²
 - ▶ “An informal, incomplete grammar of the English language runs over 1,700 pages”
- ▶ We are modelling human perception as much as we are modelling cars or numbers!

²Halevy, Alon, Peter Norvig, and Fernando Pereira. "The unreasonable effectiveness of data." IEEE Intelligent Systems 24.2 (2009): 8-12.

CIFAR 0

airplane



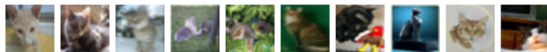
automobile



bird



cat



deer



dog



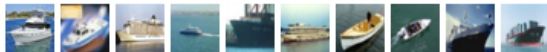
frog



horse



ship



truck



DATA AUGMENTATION

```
keras.preprocessing.image.ImageDataGenerator(featurewise_center=False,  
    samplewise_center=False,  
    featurewise_std_normalization=False,  
    samplewise_std_normalization=False,  
    zca_whitening=False,  
    rotation_range=0.,  
    width_shift_range=0.,  
    height_shift_range=0.,  
    shear_range=0.,  
    zoom_range=0.,  
    channel_shift_range=0.,  
    fill_mode='nearest',  
    cval=0.,  
    horizontal_flip=False,  
    vertical_flip=False,  
    rescale=None,  
    dim_ordering=K.image_dim_ordering())
```

CIFAR-10 DATA AUGMENTATION



KERAS CODE

```
datagen = ImageDataGenerator(  
    featurewise_center=True,  
    featurewise_std_normalization=True,  
    rotation_range=20,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    horizontal_flip=True)  
  
# compute quantities required for featurewise normalization  
# (std, mean, and principal components if ZCA whitening is applied)  
datagen.fit(X_train)  
  
# fits the model on batches with real-time data augmentation:  
model.fit_generator(datagen.flow(X_train, Y_train, batch_size=32),  
                    samples_per_epoch=len(X_train), nb_epoch=nb_epoch)
```

OUTSIDE KERAS

```
for i, (X_batch, Y_batch) in enumerate(datagen.flow(X_train, Y_train, batch_size=32)):  
    ## break once you are happy or use an incremental regressor classifier  
    ## .partial_fit
```

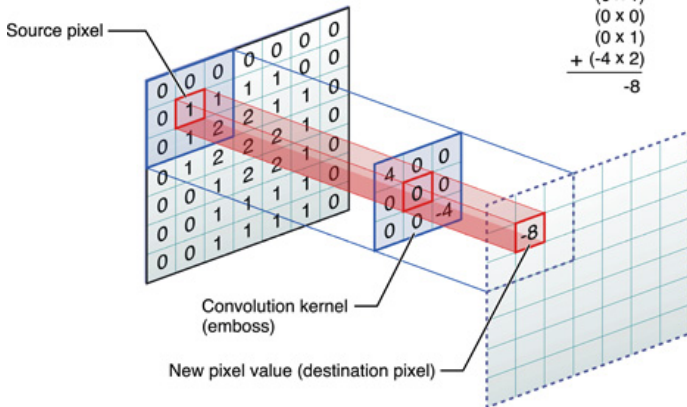
Can you do the same data augmentation operations on MNIST images?

CONVOLUTIONAL LAYERS

- ▶ Another common approach is to constraint the number of parameters
- ▶ In a layer type in neural networks become very popular due to huge successes in computer vision
- ▶ It tries to learn different filters
 - ▶ Have you ever played with photohop filters?

2D CONVOLUTIONS

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.

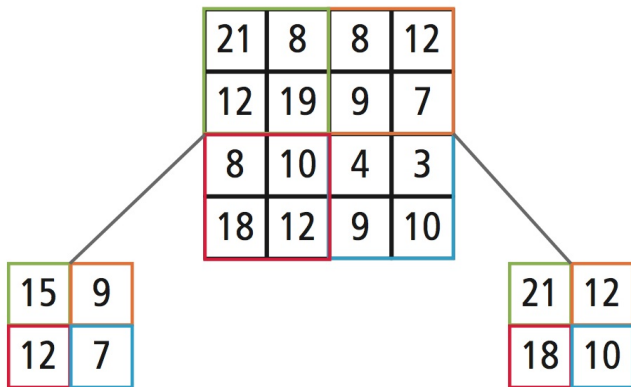


<https://developer.apple.com/library/content/documentation/Performance/Conceptual/vImage/ConvolutionOperations/ConvolutionOperations.html>

LEARNNING 2D CONVOLUTIONS

- ▶ You pass the filter over the whole image
 - ▶ Some way of treating borders
 - ▶ Padding with zeros
 - ▶ Do not calculate values if the kernel cannot fit
`http://deeplearning.stanford.edu/wiki/index.php/`
`Feature_extraction_using_convolution`
- ▶ Notice that now the size of the image doesn't matter as much
- ▶ 3x3 kernels very common

POOLING

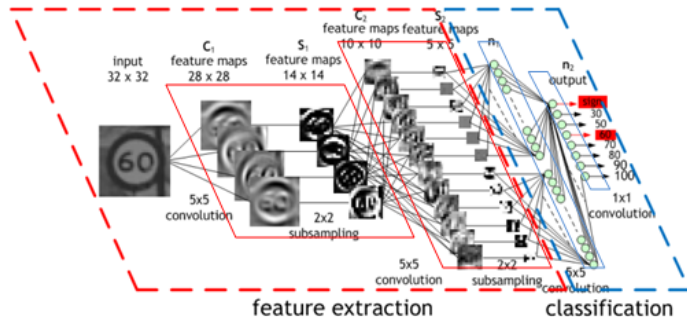


Average Pooling

Max Pooling

<http://www.embedded-vision.com/sites/default/files/technical-articles/CadenceCNN/>

VISUALISATION



<https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-core-concepts/>

All the above operations are super-optimised for GPUs

CODE

```
model = Sequential()

model.add(Convolution2D(32, 3, 3, border_mode='same',
                        input_shape=X_train.shape[1:]))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 3, 3, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))
```

VIDEO DATA

- ▶ Video is effectively a stream of images
- ▶ It has a time component
- ▶ Multiple ways of attacking this
- ▶ You can unfold and create a really large image!
- ▶ Or, 3D convolutions!
- ▶ Not too many benchmarks
 - ▶ It's hard to annotate video
<https://www.kaggle.com/c/youtube8m/>

TEXT

- ▶ There are multiple ways to treat text
- ▶ We will only see the ones here that require minimal pre-processing
 - ▶ You can create grammars, pre-process
 - ▶ These things are covered mostly under an NLP module
- ▶ We treat text as data

EMBEDDING LAYERS

- ▶ “The quick brown fox jumps over the brown lazy dog”
- ▶ Convert each word to an integer

The	1
quick	2
brown	3
fox	4
jumps	5
over	6
the	1
brown	3
lazy	7
dog	8

TRAINING (1)

- ▶ A weight matrix W is created as usual with size $(n_words, n_neurons)$
- ▶ Each row represents a word
- ▶ Each column is a specific feature/neuron
- ▶ These weights are what is passed to the follow-up layers
- ▶ What is the supervised signal?

TRAINING (2)

- ▶ Continuous bag of words
 - ▶ You are given as input n previous words and n follow up words and you try to predict the one in the middle
- ▶ $W[\textit{“Paris”}] - W[\textit{“France”}] + W[\textit{“Italy”}] \simeq W[\textit{“Rome”}]$
- ▶ $W[\textit{“king”}] - W[\textit{“man”}] \simeq W[\textit{“queen”}] - W[\textit{“woman”}]$
- ▶ You don't need to use pre-trained vectors

CODE - PREPROCESSING

```
print('Loading data...')
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=max_features)
print(len(X_train), 'train sequences')
print(len(X_test), 'test sequences')

print('Pad sequences (samples x time)')
X_train = sequence.pad_sequences(X_train, maxlen=maxlen)
X_test = sequence.pad_sequences(X_test, maxlen=maxlen)
print('X_train shape:', X_train.shape)
print('X_test shape:', X_test.shape)
```

CODE

```
model = Sequential()

model.add(Embedding(max_features,
                    embedding_dims,
                    input_length=maxlen,
                    dropout=0.2))

model.add(Convolution1D(nb_filter=nb_filter,
                        filter_length=filter_length,
                        border_mode='valid',
                        activation='relu',
                        subsample_length=1))
model.add(GlobalMaxPooling1D())

# We add a vanilla hidden layer:
model.add(Dense(hidden_dims))
model.add(Dropout(0.2))
model.add(Activation('relu'))
```

CAN WE DO IT WITHOUT NEURAL NETWORKS?

- ▶ Easy solution - create feature vector
- ▶ Very recent solution which somewhat works
 - ▶ Break the sentence/images into windowed sequences
 - ▶ i.e. generate more examples from each data point
 - ▶ Classify each of these examples
 - ▶ Combine the results into of the classifiers using a third classifier

EXAMPLE

$X[0]$ = “This film is the worst film I have ever watched. I hate the director and all the actors should be fired”

$y = 1$

- ▶ Window length of 4 (and obviously you need to turn words into numbers)

“This film is the”, 1

“film is the worst”, 1

..., 1

RECURRENT NETWORKS

- ▶ You can use convolutions to process sequences
- ▶ You can just flatten the sequence
- ▶ But often sequences have different length
 - ▶ You can pad
- ▶ How about arbitrary long sequences
 - ▶ E.g. a book?
 - ▶ Very long videos?
- ▶ Use a recurrent layer
 - ▶ Takes input of type $(n_timesteps, n_features)$

EQUATION

$$\mathbf{h}_t = (\mathbf{W} * \mathbf{h}_{t-1} + \mathbf{U} * x_t)$$

- ▶ h_{t-1} is the previous state
- ▶ \mathbf{W} is your internal weight matrix
- ▶ \mathbf{U} your external weight matrix
- ▶ x_t is the input
- ▶ Come in multiple variants - GRUs, LSTMS etc

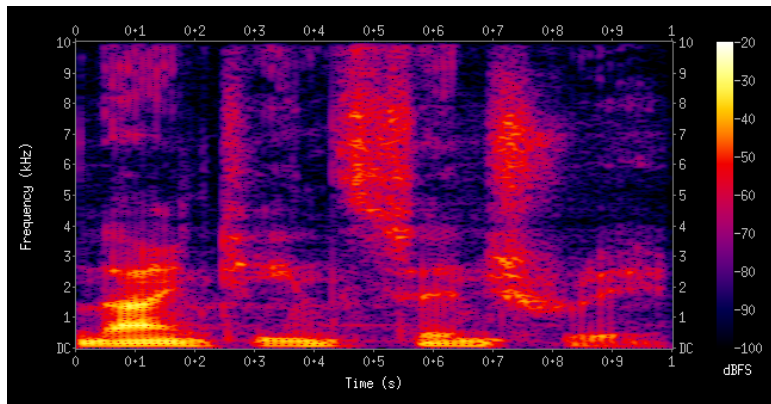
CODE

```
model = Sequential()
model.add(Embedding(max_features, 128, dropout=0.2))
model.add(LSTM(128, dropout_W=0.2, dropout_U=0.2))
model.add(Dense(1))
model.add(Activation('sigmoid'))
```

SOUND

- ▶ Sound is captured in terms of
 - ▶ Bit depth (e.g. 16bit)
 - ▶ Sample size (e.g. 44.1KHz)
- ▶ 16 bits of information are collected times a 44100 second
- ▶ You could feed this directly to an RNN/Conv network
- ▶ But usually a spectrogram is passed
 - ▶ You turn audio into an image classification problem!

SPECTROGRAM



By Aquegg - Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=5544473>
<https://yerevann.github.io/2016/06/26/combining-cnn-and-rnn-for-spoken-language-identification/>

AUTOENCODERS

- ▶ Your goal is to learn the data
- ▶ There is no other supervisory signal, but the data
- ▶ I'll give you an image as X
 - ▶ You will produce the same image as output
- ▶ Applications?
 - ▶ Image de-noising
 - ▶ Compression!

EXAMPLE CODE

```
# this is the size of our encoded representations
encoding_dim = 32

# this is our input placeholder
input_img = Input(shape=(784,))
encoded = Dense(encoding_dim, activation='relu')(input_img)
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input=input_img, output=decoded)
```

GENERATING TEXT

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- ▶ Pushed the popularity of generative methods sky-high
- ▶ Learn to generate text, given some examples

EXAMPLE

- ▶ Try to predict characters one by one
- ▶ You input a character
 - ▶ You call `.fit`
 - ▶ Network is stateful
 - ▶ i.e. it remembers where you left off!
- ▶ This way you can process super-long sequences iteratively

CODE

```
keras.layers.recurrent.Recurrent(weights=None,  
                                   return_sequences=False,  
                                   go_backwards=False,  
                                   stateful=False,  
                                   unroll=False,  
                                   consume_less='cpu',  
                                   input_dim=None,  
                                   input_length=None)
```

GANs

- ▶ Claims of being the most important advance of in AI for years
- ▶ Define a game of sorts
 - ▶ One network G generates an image/text/video
 - ▶ Another network D tries to discriminate between real and artificial examples!
 - ▶ G is trained as to produce images that D cannot differentiate!
<https://www.youtube.com/watch?v=PmC6Z0aCA0s&feature=youtu.be>

CONCLUSION

- ▶ There is more to data than just tables
- ▶ Arguably, the most interesting data is in a table format
- ▶ Again, we have just touched upon the subject
- ▶ What about sound?
 - ▶ Wavenet