

Neural networks

Spyros Samothrakis
Research Fellow, IADS
University of Essex

February 28, 2017

About

Linear function approximation with SGD

From linear regression to neural networks

Practical aspects

Conclusion

ABOUT

- ▶ We will try to get a practical understanding of neural networks
 - ▶ The topic is massive
- ▶ The field has changed names a number of times
 - ▶ Connectionist systems
 - ▶ Neural networks
 - ▶ Deep learning
- ▶ They are all the same stuff, different name
 - ▶ Re-branding

LINEAR FUNCTION APPROXIMATION

- ▶ One of the simplest ways to do a prediction
- ▶ In case we have a single variable/feature/co-variate, we are trying to learn
 - ▶ $\hat{f}_w(x) = w_1x + w_0$
- ▶ w_0 is often called the bias
- ▶ x is a sample from the data
- ▶ Let's load some data

DATA LOADING

```
df_linear = pd.DataFrame()

df_linear["AGE"] = df[["AGE"]].copy()
df_linear["INCOME"] = df[["INCOME" + str(i) for i in range(1,8)]].sum(axis = 1)

df_linear["YEARSCH"] = df[["YEARSCH"]].copy()
df_linear["ENGLISH"] = df[["ENGLISH"]].copy()
df_linear["FERTIL"] = df[["FERTIL"]].copy()
df_linear["YRSSERV"] = df[["YRSSERV"]].copy()
```

Is there another way we could have done the same?

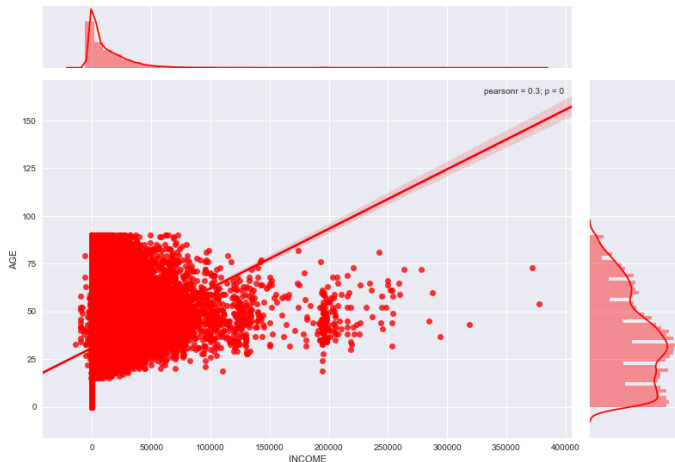
VISUALISATION

```
g = sns.jointplot("INCOME", "AGE", data=df_linear, color = "r")
```



VISUALISATION - LINEAR REGRESSION

```
g = sns.jointplot("INCOME", "AGE", data=df_linear, color = "r", kind="reg")
```



STOCHASTIC GRADIENT DESCENT (SGD)

- ▶ Let's define a cost function $J_w(x, y)$, which signifies how far we are from the solution
- ▶ We have seen such cost functions before in terms of evaluating prediction

$$\blacktriangleright J_w(x, y) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - \hat{f}_w(x^{(i)}) \right)^2$$

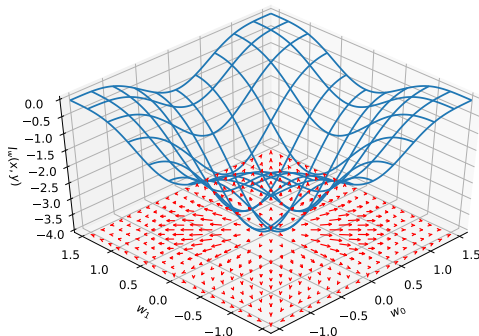
- ▶ For just one example we have:

$$\blacktriangleright J_w(x^{(0)}, y^{(0)}) = \left(y^{(0)} - \hat{f}_w(x^{(0)}) \right)^2$$

$$w = w - \eta \cdot \nabla_w J_w(x, y)$$

GRADIENTS

The gradient $\nabla_w J_w(x, y)$ is the direction of steepest descent¹



$$\nabla_w J_w(x, y) \equiv \left[\frac{\partial J_w(x, y)}{\partial w_0}, \frac{\partial J_w(x, y)}{\partial w_1}, \dots, \frac{\partial J_w(x, y)}{\partial w_n} \right]$$

¹https://commons.wikimedia.org/wiki/File:Gradient_Visual.svg

FIND THE GRADIENT

```
from sympy import Function, Symbol, latex, init_printing

y0,y1,y2 = symbols('y^(0),y^(1),y^(2)')
x0, x1, x2, x = symbols('x^(0), x^(1), x^(2), x')

f = (x*w1 + w0)

w0,w1 = symbols('w0, w1')

mse = ((f.subs(x, x0) - y0 )**2 )

mse.diff(w0)
mse.diff(w1)
```

$$\frac{\partial J_w(x,y)}{\partial w_0} = 2 \left(w_0 + w_1 x^{(0)} - y^{(0)} \right)$$
$$\frac{\partial J_w(x,y)}{\partial w_1} = 2x^{(0)} \left(w_0 + w_1 x^{(0)} - y^{(0)} \right)$$

SOME SIMPLIFICATIONS

$$[w_0, w_1] = [w_0, w_1] - \eta \cdot \left[\frac{\partial J_w(x, y)}{\partial w_0}, \frac{\partial J_w(x, y)}{\partial w_1} \right]$$

$$[w_0, w_1] = [w_0, w_1] - \eta \cdot \left[2 \left(w_0 + w_1 x^{(0)} - y^{(0)} \right), 2x^{(0)} \left(w_0 + w_1 x^{(0)} - y^{(0)} \right) \right]$$

$$[w_0, w_1] = [w_0, w_1] - \eta \cdot \left[2 \left(y^{(0)} - \hat{f}_w \left(x^{(i)} \right) \right), 2x^{(0)} \left(y^{(0)} - \hat{f}_w \left(x^{(i)} \right) \right) \right]$$

$$[w_0, w_1] = [w_0, w_1] - \eta \cdot \left[\left(y^{(0)} - \hat{f}_w \left(x^{(i)} \right) \right), x^{(0)} \left(y^{(0)} - \hat{f}_w \left(x^{(i)} \right) \right) \right]$$

MINI-BATCHES

- ▶ We used only one example
- ▶ This is called “online learning”
- ▶ But it’s common to use “mini-batches”
 - ▶ i.e., a number of samples together
- ▶ You would use every single sample
 - ▶ Often mini-batching much faster
 - ▶ Finding the right number of batches is tricky
 - ▶ Finding η is tricky

WHAT IF WE HAD THREE WEIGHTS AND TWO FEATURES?

- Let's do the derivations

SOME ANIMATIONS

<http://imgur.com/a/Hqolp>

<http://imgur.com/SmDARzn>

By Alec Reed

CREATING HIGHER ORDER FEATURES

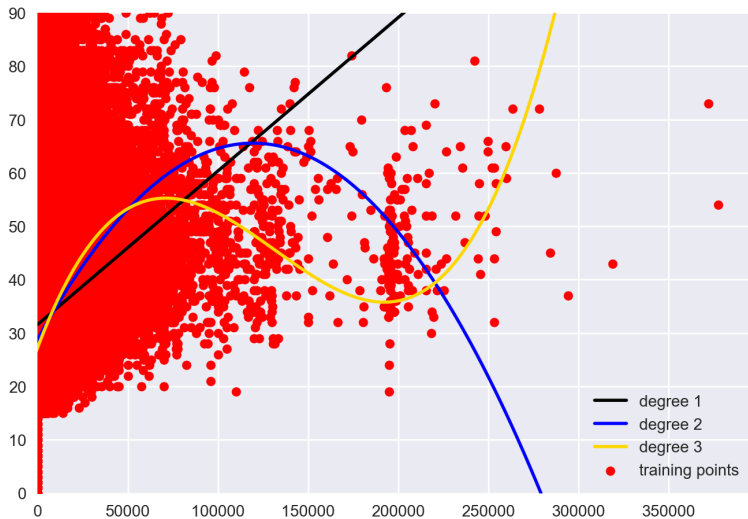
- ▶ What if the relationship between the features and the rewards is not linear?
- ▶ What if there is some other relationship?
- ▶ We can learn higher order features
- ▶ So for example, if we have x_0, x_1
 - ▶ Superscripts are rows, subscripts are columns!
- ▶ We can create new features x_0^2, x_0x_1, x_1^2

REGRESSION WITH HIGHER ORDER FEATURES

```
from sklearn.preprocessing import PolynomialFeatures

degree = 1
model = make_pipeline(PolynomialFeatures(degree), StandardScaler(), SGDRegressor())
model.fit(X, y)
y_hat = model.predict(X)
```


PLOT



ADDING ANY OTHER KIND OF FEATURE

- ▶ How about we add features in the form
 - ▶ $\max(0, x_0)$
 - ▶ $\max(0.3, x_0)$
 - ▶ $\sin(x_0)$
 - ▶ $\sin(\max(0, x_0))$
- ▶ The list is endless
- ▶ Or, even better, *discover* those features

BUILDING ABSTRACTIONS

- ▶ You rarely operate directly on the input space
 - ▶ People seem to be building abstract features
- ▶ We go from shapes and colours to forms to objects
- ▶ From primitive sounds to music
- ▶ Hierarchies of features extractors

LEARNING FEATURES

- ▶ Let's revisit our example

$$\hat{f}_w(x) = w_1 x + w_0$$

- ▶ Let's add some depth using another function

$$\hat{f}_{w^{(0)}}(x) = w_1^{(0)} \hat{g}_w(x) + w_0^{(0)}$$

$$\hat{g}_{w^{(1)}}(x) = \max(w_1^{(1)} x + w_0^{(1)}, 0.3)$$

- ▶ The neuron above (g) are known as a “leaky rectifier” neuron
 - ▶ But it's just a function - the neuron terminology comes from the original inspiration
 - ▶ Can you see why the name “deep learning” stuck?

LEARNING

- Use the same error function to find errors

$$J_w(x^{(0)}, y^{(0)}) = \left(y^{(0)} - \hat{f}_w(x^{(i)})\right)^2$$

- Update each weight using the same method

$$w = w - \eta \cdot \nabla_w J_w(x, y)$$

$$\nabla_w J_w(x, y) \equiv \left[\frac{\partial J_w(x, y)}{\partial w_0}, \frac{\partial J_w(x, y)}{\partial w_1}, \dots, \frac{\partial J_w(x, y)}{\partial w_n} \right]$$

- Use the chain rule to calculate partial derivatives

$$\frac{d}{dx} [f(g(x))] = f'(g(x))g'(x)$$

CONVEXITY

- ▶ We are going downhill
- ▶ What if the error surface is smooth, it's all good
- ▶ But what if the error surface is rugged mess with ups and downs?
- ▶ Local optima are not global optima
- ▶ Thus, NN training often somewhat unstable

THEANO, TENSORFLOW

- ▶ As you understand doing all these calculations by hand is tedious
- ▶ Every new type of neuron or layer you invent would require a new gradient calculation
- ▶ Auto-differentiation
 - ▶ Theano
 - ▶ Tensorflow
- ▶ You just describe your network/graph
- ▶ The rest is done automatically by the machine

BUILDING A NETWORK

- ▶ There are multiple packages for neural networks
- ▶ I think the most popular is Keras
- ▶ Though a massive number of people are using Lasagne

KERAS API

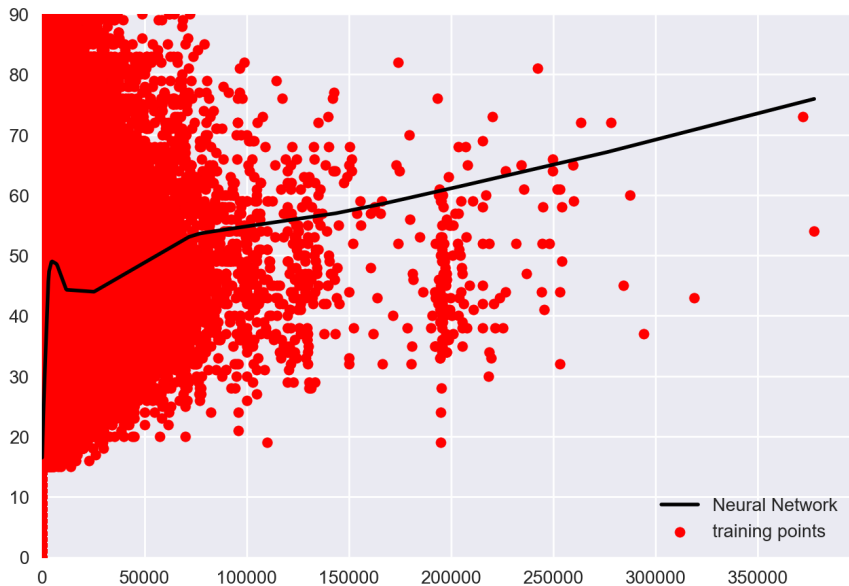
```
inputs = Input(shape=(1,))

x = Dense(64)(inputs)
x = PReLU()(x) # Non-linearity
x = Dense(64)(x)
x = PReLU()(x) # Non-linearity
predictions = Dense(1, activation='linear')(x)

model = Model(input=inputs, output=predictions)
model.compile(optimizer='adam', loss='mse')

model.fit(...)
```

OUTPUT PLOT



INPUT LAYER

```
inputs = Input(shape=(22,))
```

- ▶ The first layer
- ▶ Does not contain any neurons
 - ▶ Terminology changes between authors
- ▶ “Number of columns”
- ▶ “Number of features”

HIDDEN LAYERS

```
x = PReLU()(x) # Non-linearity  
x = Dense(64)(x)  
x = PReLU()(x) # Non-linearity
```

- ▶ “Dense” layers specify the number of neurons
- ▶ You need to follow them up with a non-linear transformation
- ▶ Alternately you can pass this as parameter “activation”
- ▶ Non-linearity is often called the “activation function”

OUTPUT LAYER

```
predictions = Dense(1, activation='linear')(x)
model = Model(input=inputs, output=predictions)
```

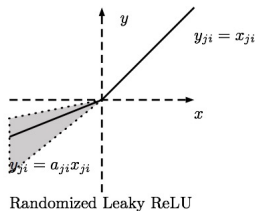
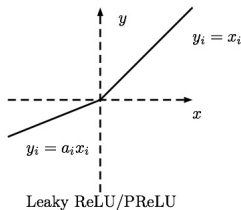
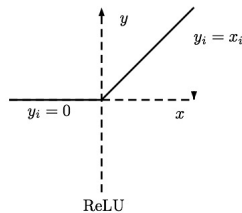
- ▶ The output layer
- ▶ A normal Dense layer
- ▶ “Number of outputs”
- ▶ You can predict multiple outputs at once

COST (OR OBJECTIVE) FUNCTIONS

- ▶ You can do both regression and classification with neural networks
- ▶ All it changes is the output layer and the objective function
- ▶ Categorical cross-entropy for classification
 - ▶ You need to convert your outputs to one-hot encoding
- ▶ Mean squared error for prediction

NEURON TYPES

- Also called activation functions or non-linearities



²<http://lamda.nju.edu.cn/weixs/project/CNNTricks/imgs/relufamily.png>

LEARNING ALGORITHMS

- ▶ We have just seen SGD, which forms the basis of every modern learning algorithm for NNs
 - ▶ It will be dethroned eventually, but nobody knows from what
- ▶ SGD does not take into account specific information about individual weights
 - ▶ Not the curvature of the search space
- ▶ There are methods that auto-adapt the learning rate per individual weight
 - ▶ Adam
 - ▶ RMSProp
- ▶ Adam should be your default option

REGULARISATION

- ▶ Overfitting is a massive problem - with enough variables you can fit anything
- ▶ ℓ_1 : Turn your weights down to zero
- ▶ ℓ_2 : Make all your weights small
- ▶ You can define the regularisation strength
- ▶ Dropout:
 - ▶ Remove some of the weights during training uniform random
 - ▶ During testing/using put use everything, but toned down
 - ▶ Implemented as a distinct layer

(BATCH) NORMALISATION

- ▶ Implemented as a different layer
- ▶ Covariate shift impact learning negatively
- ▶ Normalise the weights of each layer

$$\hat{w} = \frac{w - \mu}{\sigma}$$

$$y = \gamma \hat{w} + \beta$$

- ▶ γ and β are special types of parameters, to be learned as well
 - ▶ They scale and shift a bit
- ▶ Stops “internal covariate shift”
 - ▶ Accelerates learning

WEIGHT SHARING / MULTIPLE OUTPUTS

- ▶ You can create pathways for different things that share the same weights
- ▶ Thus imposing stronger regularisation
 - ▶ Use same features for different tasks
- ▶ Very effective method of generalising
- ▶ You can use different pathways within the same network

METRICS/ CALLBACKS

- ▶ When should you stop training?
 - ▶ After i iterations over the whole dataset
 - ▶ After the validation error stops improving
- ▶ Save best validation model
 - ▶ ... and use this for testing
- ▶ Reduce learning rates, log etc
- ▶ “Callbacks”

CREATING CUSTOM LAYERS

3

```

from keras import backend as K
from keras.engine.topology import Layer
import numpy as np

class MyLayer(Layer):
    def __init__(self, output_dim, **kwargs):
        self.output_dim = output_dim
        super(MyLayer, self).__init__(**kwargs)

    def build(self, input_shape):
        # Create a trainable weight variable for this layer.
        self.W = self.add_weight(shape=(input_shape[1], self.output_dim),
                                initializer='random_uniform',
                                trainable=True)
        super(MyLayer, self).build() # Be sure to call this somewhere!

    def call(self, x, mask=None):
        return K.dot(x, self.W)

    def get_output_shape_for(self, input_shape):
        return (input_shape[0], self.output_dim)

```

³<https://keras.io/layers/writing-your-own-keras-layers/>

CONCLUSION

- ▶ We have touched upon neural networks
- ▶ It's a really hot topic right now
- ▶ Requires a bit of dedication
- ▶ New algorithms are coming out every day